

Apache Streaming

Why Spark Streaming

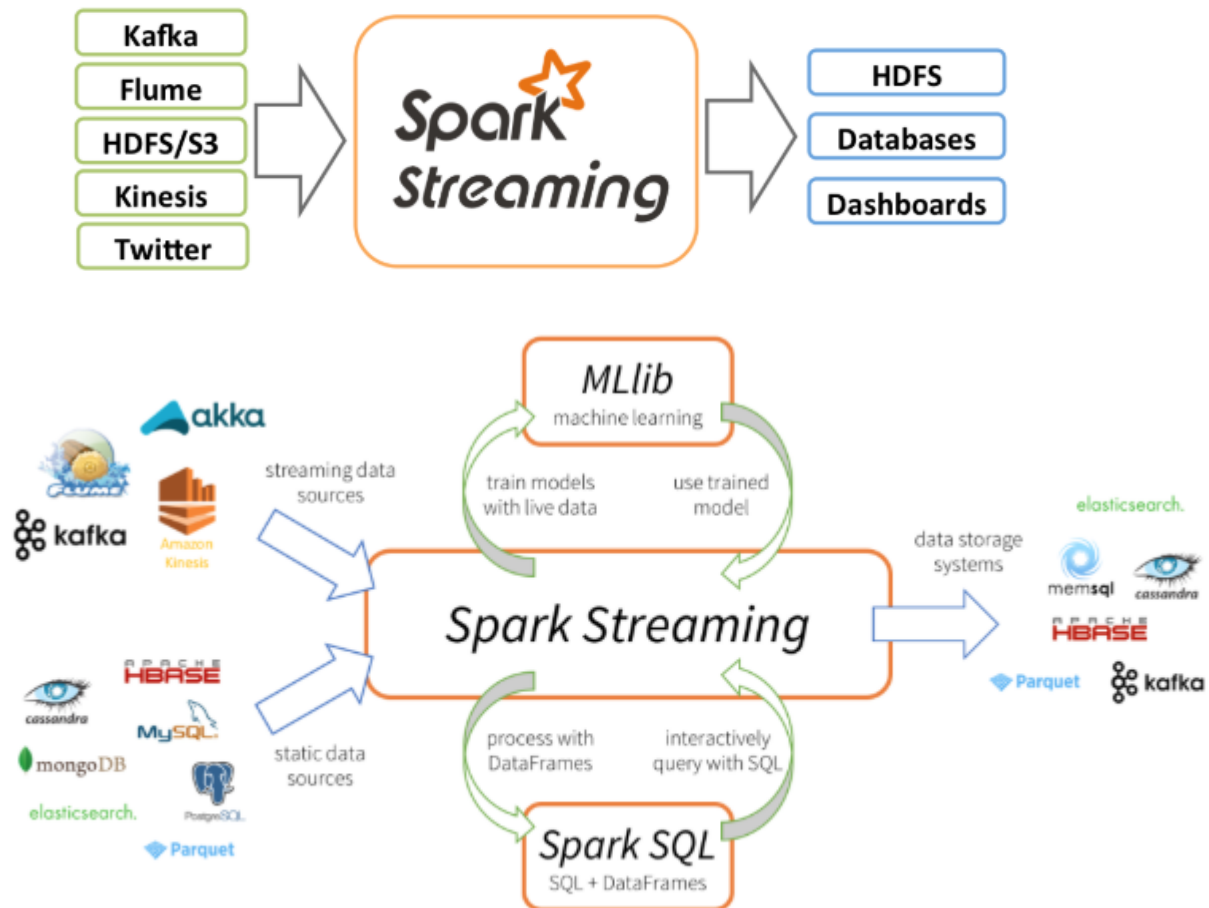
- ▶ Many big-data applications need to process large data streams in real time
- ▶ Streaming is a technique for transferring data so that it can be processed as a steady and continuous stream.
- ▶ Apache Spark has provided an unified engine that natively supports both batch and streaming workloads.



Spark Streaming is used to stream real-time data from various sources like Twitter, Stock Market and Geographical Systems and perform powerful analytics to help businesses.

Spark Streaming

- ▶ Spark Streaming is an extension of the core Spark API that enables scalable, high-throughput, fault-tolerant stream processing of live data streams.



Basic sources:

- ▶ Sources directly available in the StreamingContext API. Examples: file systems, and socket connections.

Advanced sources:

- ▶ Sources like Kafka, Flume, Kinesis, etc. are available through extra utility classes.

DStream

- ▶ Discretized Stream or DStream is the basic abstraction provided by Spark Streaming
- ▶ Represents continuous stream of data
- ▶ Implemented as a sequence of RDDs
- ▶ Created from streaming input sources
- ▶ Created by applying transformations on existing DStreams



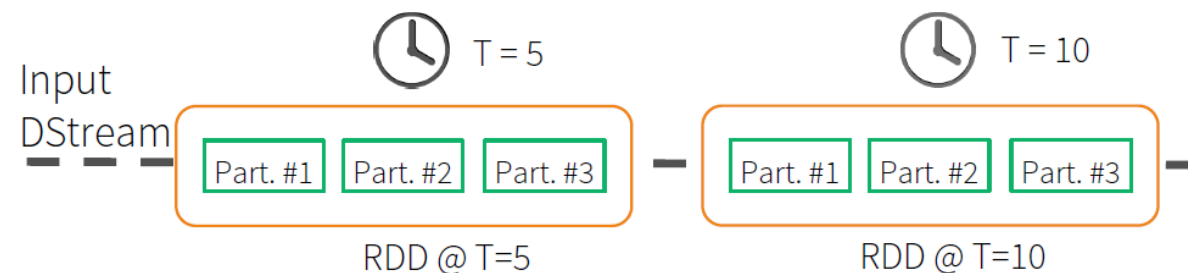
Block Interval:

- ▶ Controls how often we create partition
- ▶ 200ms by default

Batch Interval:

- ▶ Timer till it needs to be buffered for creating RDD
- ▶ It is programmatically controlled while creating Dstream

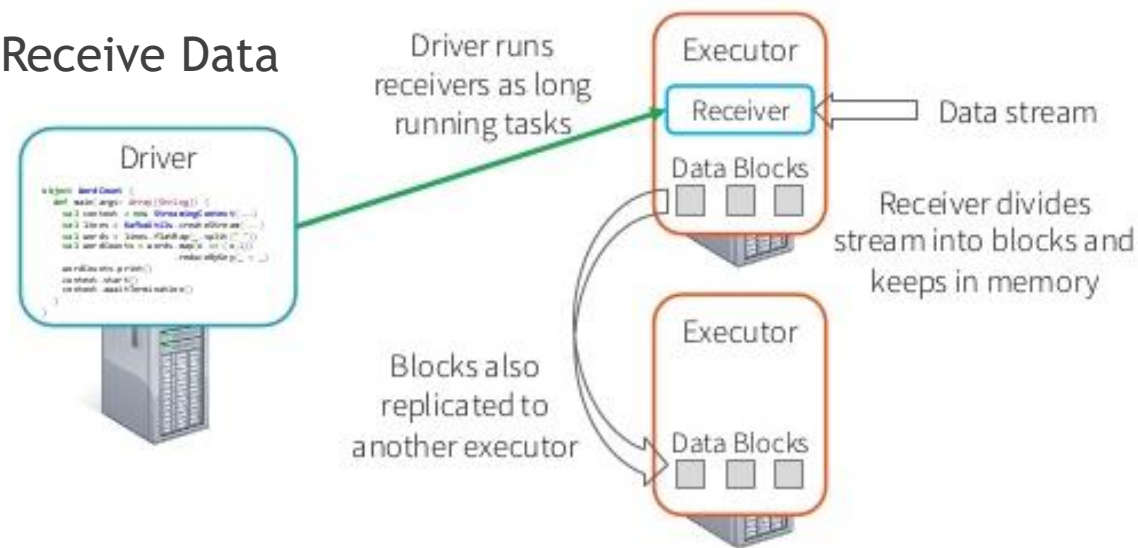
Block Interval = 200 ms
Batch interval = 5 seconds




One Partition is created every 200 ms
One RDD is created every 5 seconds

Spark Streaming - Workflow

Receive Data



Process Data



`val context = new StreamingContext(conf, Seconds(1))`

`val lines = context.socketTextStream(...)`

DStream: represents a data stream

`val words = lines.flatMap(_.split(" "))`

Transformations: transform data to create new DStreams

`val wordCounts = words.map(x => (x, 1)).reduceByKey(_+_)`

`wordCounts.print()`

Print the DStream contents on screen

`context.start()`

Start the streaming job

Receiver-based:

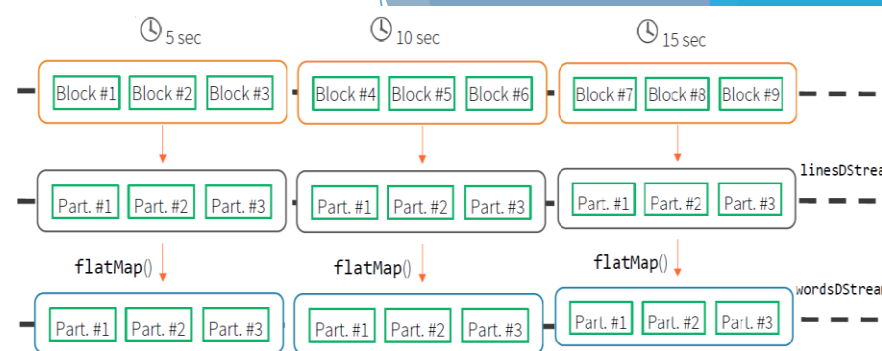

- ▶ For sources without a buffer, like TCP sockets
- ▶ A long running task called the Receiver is launched in one of the executors

Receiver-less or Direct:

- ▶ For sources with a buffer, like Kafka
- ▶ Streaming context gets partition info from source and then creates RDD

Transformations

Transformations modify data from one DStream to another



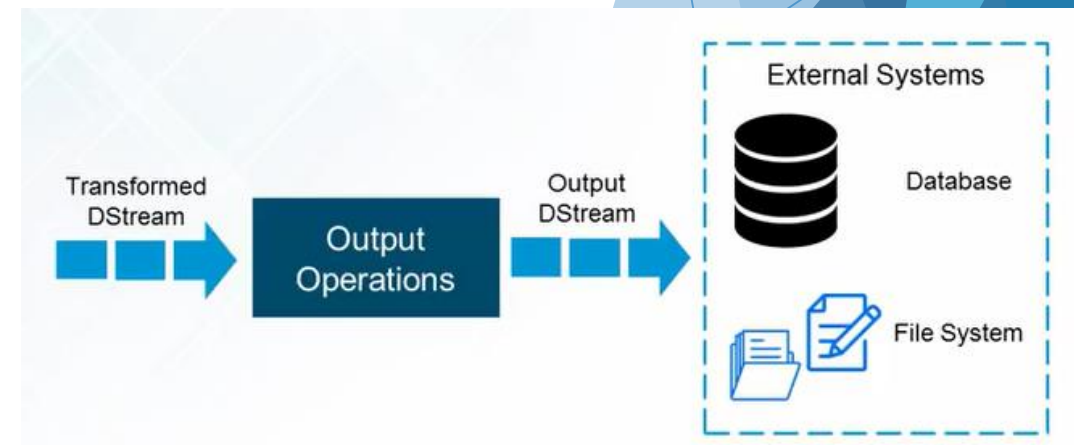
Transformations	Description
<ul style="list-style-type: none">map(func)flatMap(func)filter(func)repartition(numPartitions)union(otherStream)count()reduce(func)reduceByKey(func, [numTasks])countByValue()join(otherStream, [numTasks])cogroup(otherStream, [numTasks])	<ul style="list-style-type: none">All these return new DStream
transform(func)	<ul style="list-style-type: none">It can be used to apply any RDD operation that is not exposed in the DStream API

Output Operations

Actions	Description
<ul style="list-style-type: none">• <code>saveAsTextFiles(...)</code>• <code>saveAsObjectFiles(...)</code>• <code>saveAsHadoopFile(...)</code>	<ul style="list-style-type: none">• Save DStream's contents as text files/ SequenceFiles /Hadoop files
<code>print()</code>	<ul style="list-style-type: none">• Prints the first ten elements of every batch of data in a DStream• This is useful for development and debugging.
<code>foreachRDD(func)</code>	<ul style="list-style-type: none">• To call an action that is available in RDD API not in Dstream API

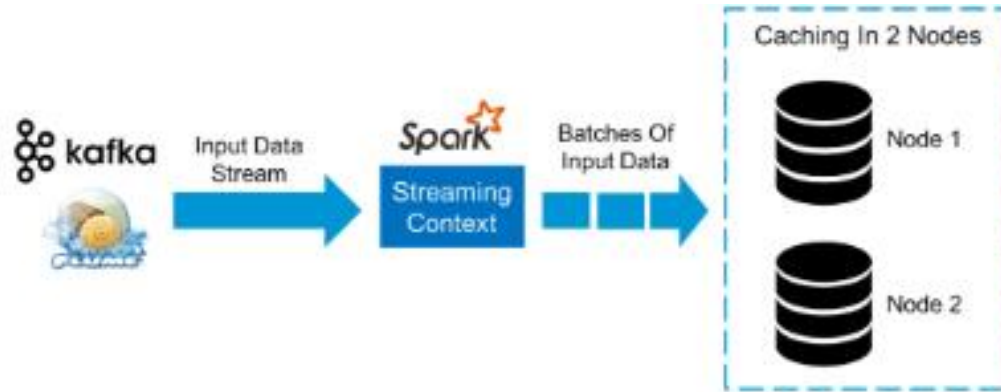
Power Combo = Dstream + RDD

- ▶ Combine live data streams with historical data
- ▶ Combine streaming with MLlib, GraphX algos
- ▶ Query streaming data using SQL



Caching

- ▶ DStreams allow developers to cache/ persist the stream's data in memory.
- ▶ This is useful if the data in the DStream will be computed multiple times.
- ▶ This can be done using the `persist()` method on a DStream.



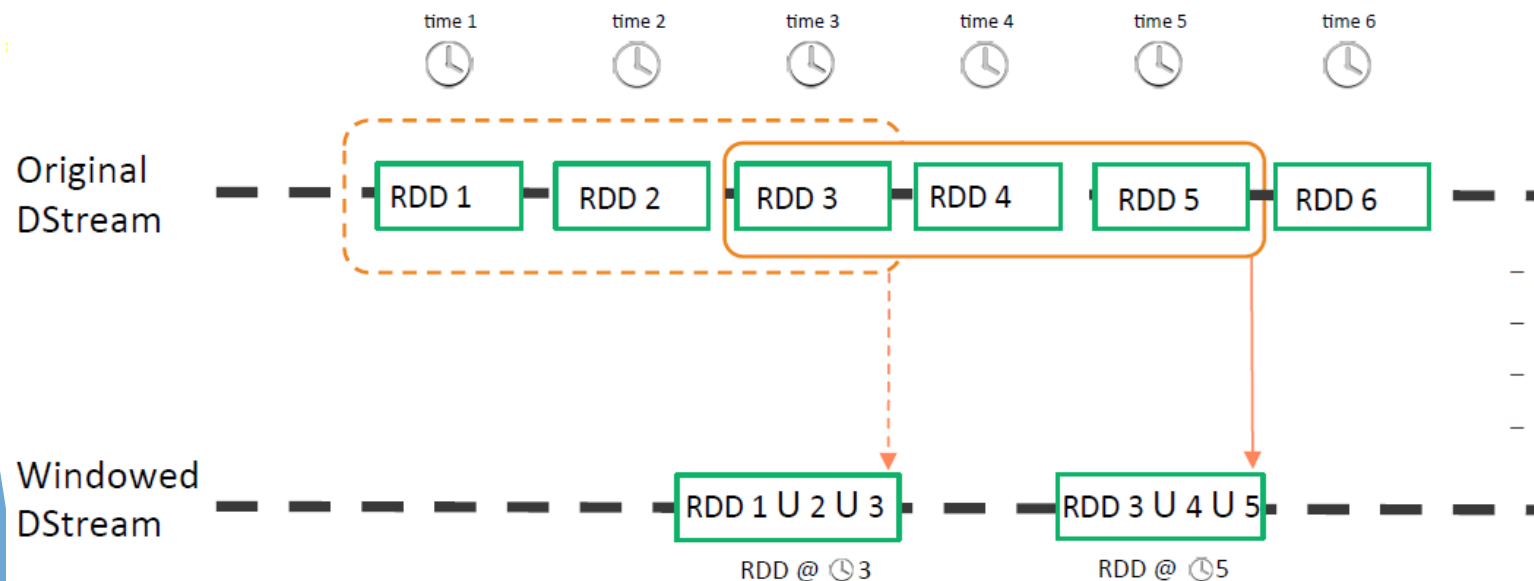
Program Structure

- ▶ StreamingContext is the main entry point for Spark Streaming functionality
 - `socketTextStream(host, port)`
 - `fileStream(directory)`, `textFileStream(directory)`
 - `actorStream(actorProps, actorName)`
 - `queueStream(queueOfRDDs)`
- ▶ Specify the batchInterval in the program
- ▶ Create Dstream object by specifying source
- ▶ Construct DStream DAG
- ▶ Once the streaming process is started the DAG cannot be changed
- ▶ Once streaming context is stopped, streaming context will be garbage collected
- ▶ Stopping StreamingContext stops SparkContext as well by default
- ▶ `StreamingContext.start()` to kickoff the streaming process
- ▶ `StreamingContext.awaitTermination()` to wait for termination
- ▶ `StreamingContext.stop()` to stop the streaming process



Window Operations

- ▶ Apply transformations over a sliding window of data
- ▶ Every time the window slides over a source DStream, the source RDDs that fall within the window are combined and operated to produce the RDDs
- ▶ Window length - The duration of the window
- ▶ Sliding interval - The interval at which the window will slide or move forward
- ▶ Ex: Window Length = 30 Secs ; Sliding Interval = 20 Secs



- `window(windowLength, slideInterval)`
- `countByWindow(windowLength, slideInterval)`
- `reduceByWindow(func, windowLength, slideInterval)`
- `reduceByKeyAndWindow(...)`

Fault Tolerance

Executor/Receiver failure:

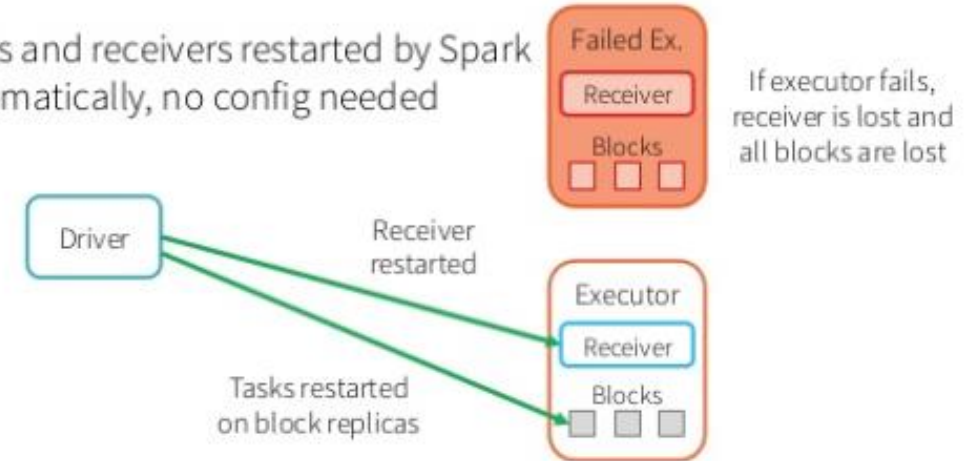
- ▶ Batches of input data are replicated in memory for fault-tolerance
- ▶ Data lost due to worker failure, can be recomputed from replicated input data
- ▶ Data buffered but not yet replicated will be lost. This can be handled by WAL
- ▶ WAL - Write Ahead Log
- ▶ Synchronously save the received data to a fault tolerant storage

Driver failure:

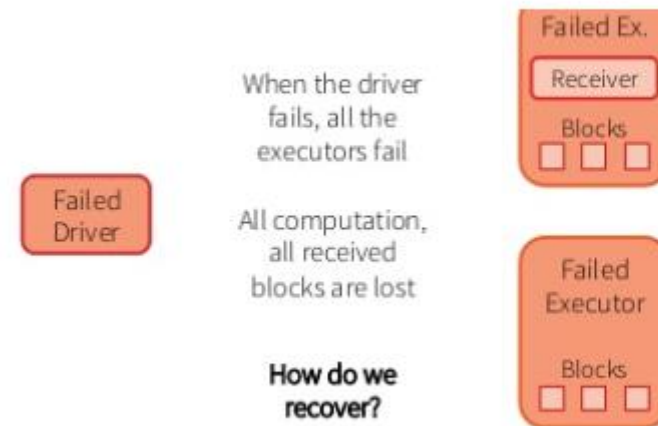
- ▶ Executors are lost including any Receivers they were running

What if an executor fails?

Tasks and receivers restarted by Spark automatically, no config needed



What if the driver fails?






Checkpoint - Driver Failure

Recovering Driver with Checkpointing

DStream Checkpointing:

Periodically save the DAG of DStreams to fault-tolerant storage

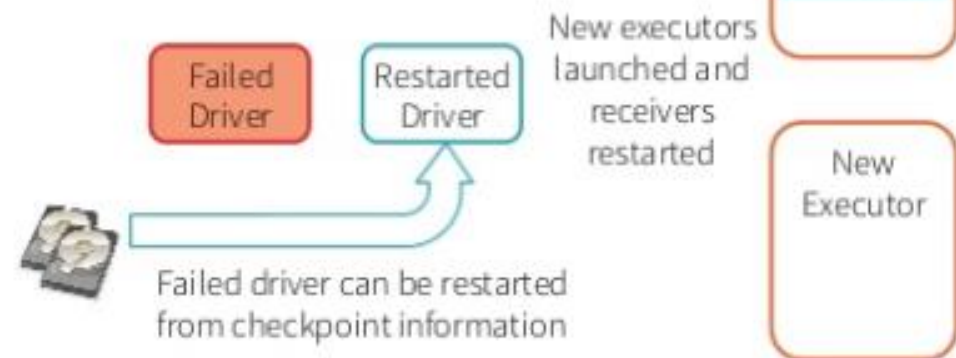


Metadata Checkpointing

Recovering Driver w/ DStream Checkpointing

DStream Checkpointing:

Periodically save the DAG of DStreams to fault-tolerant storage



Checkpointing

- ▶ A process of writing received records at checkpoint intervals to HDFS is checkpointing
- ▶ Checkpointing creates fault-tolerant stream processing pipelines.
- ▶ DStreams can checkpoint input data at specified time intervals.

Metadata Checkpointing:

- ▶ Saving of the information defining the streaming computation to fault-tolerant storage like HDFS.
- ▶ This is used to recover from failure of the node running the driver of the streaming application

Data Checkpointing:

- ▶ Saving of the generated RDDs to reliable storage
- ▶ This is necessary in some stateful transformations that combine data across multiple batches.

```
val checkpointDir = ...

def creatingFunc(): StreamingContext = {
    val newSsc = ... // create and setup a new StreamingContext

    newSsc.checkpoint(checkpointDir) // enable checkpointing


    ...
}

// Recreate context from checkpoints info in checkpointDir, or create a new one by
val ssc = StreamingContext.getOrCreate(checkpointDir, creatingFunc _)
```

Structured Streaming

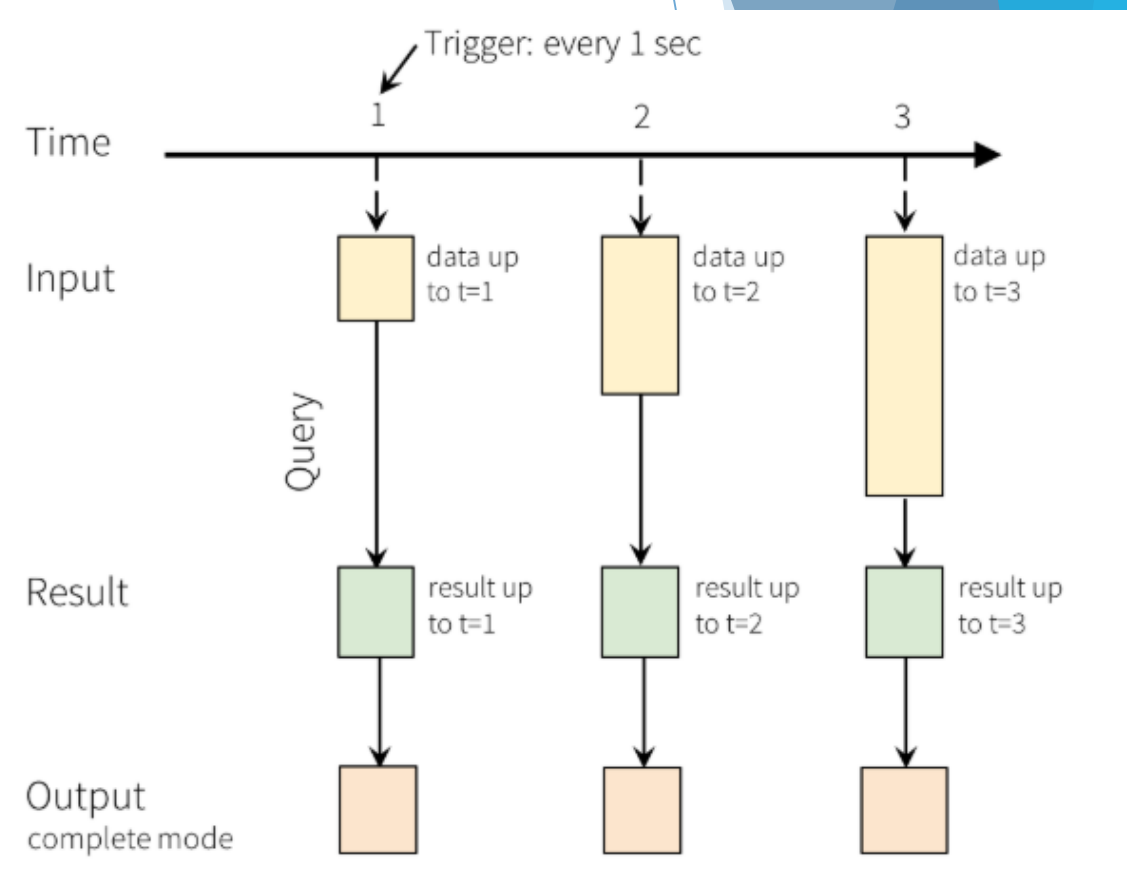
Structured Streaming

- ▶ Structured Streaming is a stream processing engine built on the Spark SQL engine
- ▶ Streaming computation can be done on the same way as that of a batch computation on static data.
- ▶ Rich, unified and high level APIs (Dataframes & Dataset)
- ▶ No more dealing with RDD directly!
- ▶ The biggest difference is latency and message delivery guarantees
- ▶ Treat stream of data as unbounded tables
- ▶ Both Structured Streaming and Streaming with DStreams use micro-batching



Processing Model

- ▶ Structured Streaming treats all the data arriving as an unbounded input table.
- ▶ Each new item in the stream is like a row appended to the input table.
- ▶ The developer then defines a query on this input table, as if it were a static table
- ▶ Spark automatically converts this batch-like query to a streaming execution plan (*incrementalization*)
- ▶ A query on the input will generate the “Result Table” which will be written to an output sink
- ▶ Developers specify triggers to control when to update the results.
- ▶ Each time a trigger fires, Spark checks for new data (new row in the input table), and incrementally updates the result.



RDD API Streaming vs Structured Streaming



Apache Spark API



Spark Streaming (D Streams) – RDD API

- RDD based Streaming
- Batch time has to be specified inside the application
- Higher latency is batch size is not properly assigned
- Triggers are bound often not customizable and are bound by batch and window time

```
val ssc= new StreamingContext(conf, Seconds(1))
val lines=ssc.socketTextStream("localhost",9999)
val words=lines.flatMap(_.split(" "))
val pairs=words.map(word=>(word,1))
val wordCount=pairs.reduceByKey(_+_).wordCounts.print()
ssc.start()
```

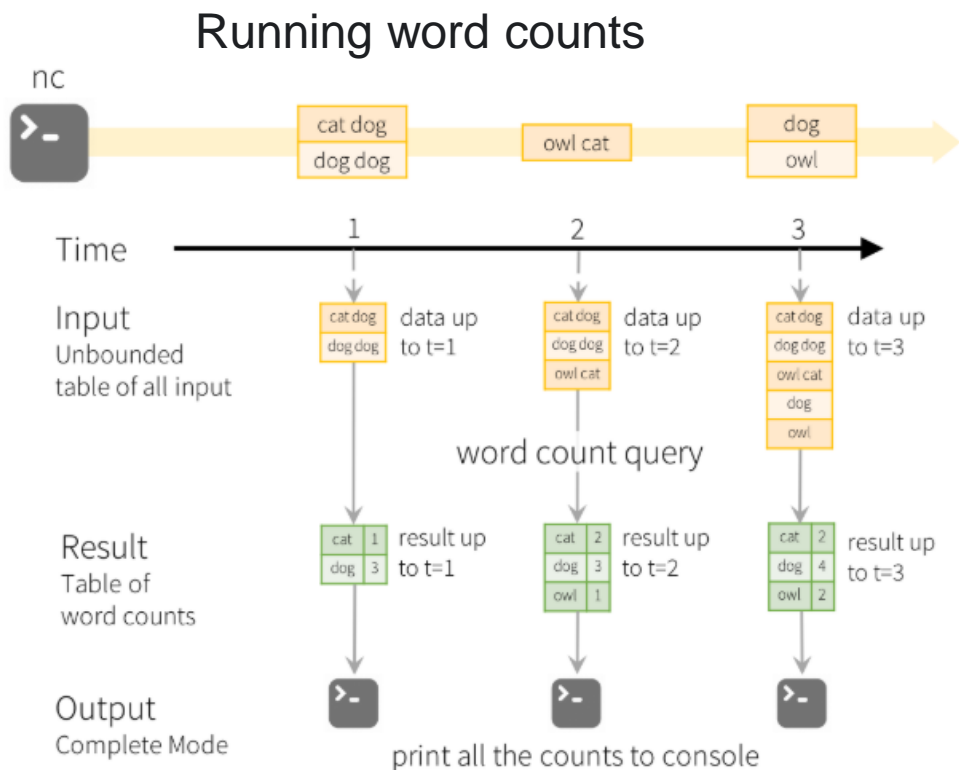
Spark Structured Streaming (DataFrame/Dataset API)

- Dataframe/Dataset based Streaming
- Uses micro-batch
- Low latency achieved through micro-batches
- Highly customizable triggers that can manipulate processing and event times separately

```
val lines=spark.readStream.format("socket").option("host","localhost").option("port","9999").load()
val words=lines.as[String].flatMap(_.split(" "))
val wordCounts=words.groupby("value").count()
val query=wordCounts.writeStream.outputMode("complete").format("console").start()
```

Output Modes

- ▶ **Append:** Only the new rows appended to the result table since the last trigger will be written to the external storage.
- ▶ **Complete:** The entire updated result table will be written to external storage
- ▶ **Update:** Only the rows that were updated in the result table since the last trigger will be changed in the external storage.



```
val lines = spark.readStream
  .format("socket")
  .option("host", "localhost")
  .option("port", 9999)
  .load()

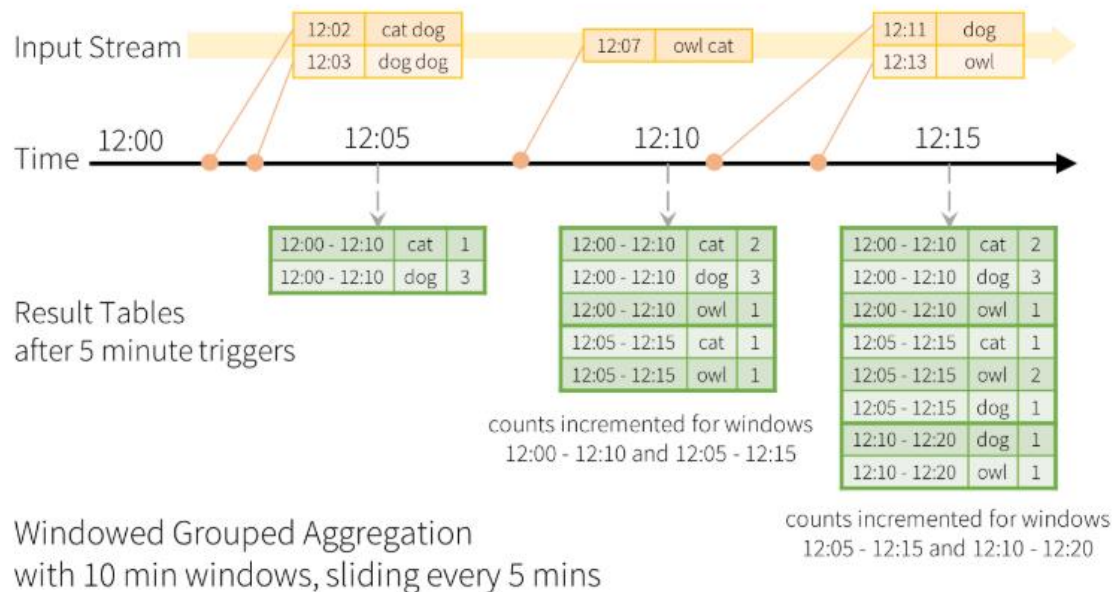
// Split the lines into words
val words = lines.as[String].flatMap(_.split(" "))

// Generate running word count
val wordCounts = words.groupBy("value").count()

val query = wordCounts.writeStream
  .outputMode("complete")
  .format("console")
  .start()
```

Window Operations

- ▶ Aggregations over a sliding event-time window are straightforward with Structured Streaming
- ▶ We want to count words within 10 minute windows (*similar to window length*), updating every 5 minutes (*similar to sliding interval*)
- ▶ That is, word counts in words received between 10 minute windows 12:00 - 12:10, 12:05 - 12:15, 12:10 - 12:20, etc
- ▶ consider a word that was received at 12:07. This word should increment the counts corresponding to two windows 12:00 - 12:10 and 12:05 - 12:15



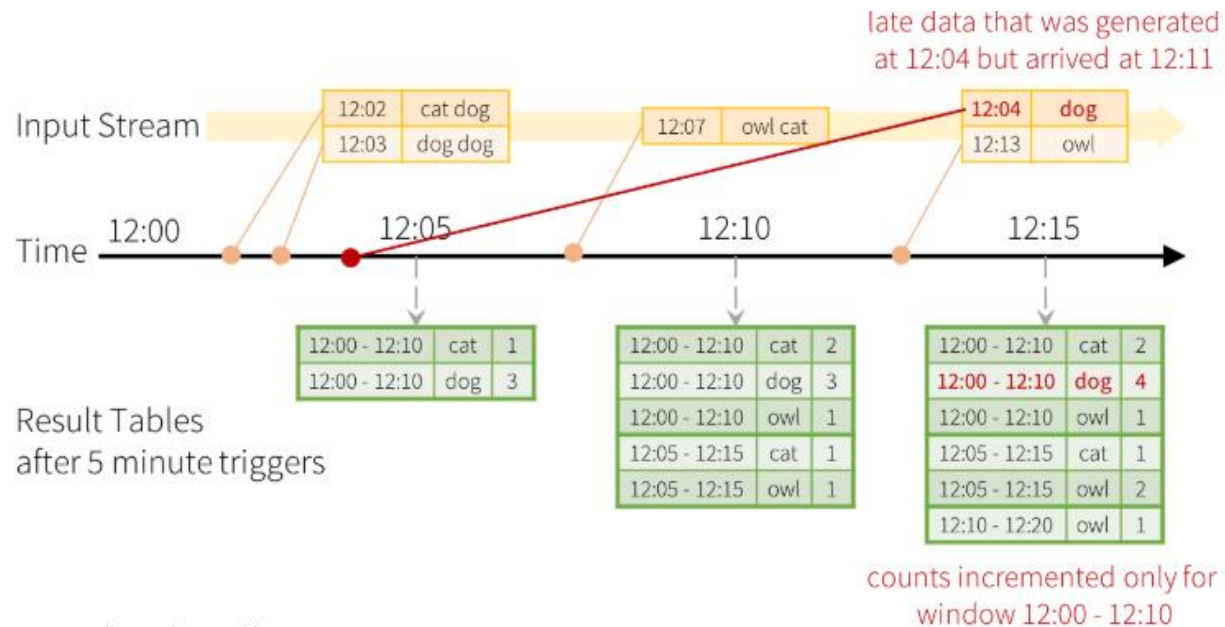
```
import spark.implicits._

val words = ... // streaming DataFrame of schema { timestamp: Timestamp, word: String }

// Group the data by window and word and compute the count of each group
val windowedCounts = words.groupBy(
  window($"timestamp", "10 minutes", "5 minutes"),
  $"word"
).count()
```

Late Data and Watermarking

- ▶ Now consider what happens if one of the events arrives late to the application. For example, say, a word generated at 12:04 (i.e. event time) could be received by the application at 12:11
- ▶ The application should use the time 12:04 instead of 12:11 to update the older counts for the window 12:00 - 12:10
- ▶ Structured Streaming can maintain the intermediate state for partial aggregates for a long period of time such that late data can update aggregates of old windows correctly



```
import spark.implicits._

val words = ... // streaming DataFrame of schema { timestamp: Timestamp, word: String }

// Group the data by window and word and compute the count of each group
val windowedCounts = words
  .withWatermark("timestamp", "10 minutes")
  .groupBy(
    window($"timestamp", "10 minutes", "5 minutes"),
    $"word")
  .count()
```

Output Sinks

Sink	Supported Output Modes	Options
File Sink	Append	<p>path: path to the output directory, must be specified.</p> <p>For file-format-specific options, see the related methods in DataFrameWriter (Scala/Java/Python/R). E.g. for "parquet" format options see <code>DataFrameWriter.parquet()</code></p>
Kafka Sink	Append, Update, Complete	See the Kafka Integration Guide
Foreach Sink	Append, Update, Complete	None
ForeachBatch Sink	Append, Update, Complete	None
Console Sink	Append, Update, Complete	<p>numRows: Number of rows to print every trigger (default: 20)</p> <p>truncate: Whether to truncate the output if too long (default: true)</p>
Memory Sink	Append, Complete	None

```
writeStream
  .format("parquet")           // can be "orc", "json", "csv", etc.
  .option("path", "path/to/destination/dir")
  .start()
```

```
writeStream
  .format("kafka")
  .option("kafka.bootstrap.servers", "host1:port1,host2:port2")
  .option("topic", "updates")
  .start()
```

```
writeStream
  .format("console")
  .start()
```

```
writeStream
  .format("memory")
  .queryName("tableName")
  .start()
```


Triggers

- ▶ The trigger settings of a streaming query defines the timing of streaming data processing, whether the query is going to be executed as micro-batch query with a fixed batch interval or as a continuous processing query

Once

- ▶ When you set the trigger option to “Once” it processes only once and then terminates the stream.

Processing Time

- ▶ This is the most widely used and recommended practice in Spark Structured Streaming.
- ▶ The trigger option of processing time gives you better control over how often micro batch jobs should get triggered
- ▶ For example, if you set a trigger interval of “20 seconds” then for every 20 seconds a micro batch will process a batch of records

Continuous

- ▶ As of Spark 2.4.0, it’s still experimental. It enables you to process records in milliseconds which would otherwise take seconds in micro batch.
- ▶ Long-running task is created per write stream and processed as quickly as possible

```
import org.apache.spark.sql.streaming.Trigger

// Default trigger (runs micro-batch as soon as it can)
df.writeStream
  .format("console")
  .start()

// ProcessingTime trigger with two-seconds micro-batch interval
df.writeStream
  .format("console")
  .trigger(Trigger.ProcessingTime("2 seconds"))
  .start()

// One-time trigger
df.writeStream
  .format("console")
  .trigger(Trigger.Once())
  .start()
```

```
spark.readStream.format("kafka")
  .option("kafka.bootstrap.servers", "...")
  .option("subscribe", "topic")
  .load()
  .selectExpr("cast (value as string) as json")
  .select(from_json("json", schema).as("data"))
  .writeStream .format("parquet")
  .option("path", "/parquetTable/")
  .trigger("1 second")
  .option("checkpointLocation", "...")
  .start()
```

Trigger: when to process data -
Fixed interval micro-batches - As fast as possible micro-batches -
Continuously
Checkpoint location: for tracking

Managing Streaming Queries

- ▶ The StreamingQuery object created when a query is started can be used to monitor and manage the query.

```
val query = df.writeStream.format("console").start()  // get the query object

query.id          // get the unique identifier of the running query that persists across restarts from checkpoint data

query.runId       // get the unique id of this run of the query, which will be generated at every start/restart

query.name        // get the name of the auto-generated or user-specified name

query.explain()   // print detailed explanations of the query

query.stop()      // stop the query

query.awaitTermination() // block until query is terminated, with stop() or with error

query.exception   // the exception if the query has been terminated with error

query.recentProgress // an array of the most recent progress updates for this query

query.lastProgress // the most recent progress update of this streaming query
```

Reference

- ▶ Spark Streaming Guide

<https://spark.apache.org/docs/0.9.2/streaming-programming-guide.html>

- ▶ Streaming Example - Kafka

<http://spark.apache.org/docs/latest/streaming-kafka-integration.html>

- ▶ Structured Streaming

<https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>



Thank You

Keerthiga Barathan